Jane Castleman

Email: castlej@stanford.edu
URL: https://jane2620.github.io/

Researcher passionate about measuring the social impacts of AI, drawing upon theory from machine learning, statistics, measurement science, and consumer protection law. Extensive experience interacting with real users to execute end-to-end research studying deployed systems.

Education

2026-Curr **Stanford University**, PhD in Computer Science

2024-2026 **Princeton University**, MSE in Computer Science (w/ thesis) GPA: 4.0 / 4.0
Advisor: Prof. Aleksandra Korolova
Relevant Graduate Courses: Theoretical ML; Deep Learning Theory; AI Safety & Alignment; Security & Privacy; Deep Dive into LLMs; CS, Law, & Policy; Fairness in ML.

2020-2024 **Princeton University**, BSE in Computer Science GPA: 3.9 / 4.0
Advisors: Prof. Katrina Ligett & Prof. Aleksandra Korolova Magna cum laude
Minors in Statistics & Machine Learning, Technology Policy

Publications & Talks

Conference papers

2026 "Measuring Validity and Fairness in LLM-based Resume Screening"
Jane Castleman, Zeyu Shen, Blossom Metevier, Max Springer, & Aleksandra Korolova, 2026.
Second Conference of the International Association of Safe & Ethical AI (IASEAI '26). [Paper]

"The Geometry of Alignment Collapse: When Fine-Tuning Breaks Safety"
Max Springer, Chung Peng Lee, Bohdan Turbal, Blossom Metevier, Hayoung Jung, **Jane Castleman**, Zeyu Shen, Aleksandra Korolova, 2026.
Fortieth Annual Conference on Neural Information Processing Systems (NeurIPS '26). [Paper]

2025 "Adultification Bias in LLMs and T2I Models"
Jane Castleman and Aleksandra Korolova, 2025.
ACM Conference on Fairness, Accountability, and Transparency (FAccT '25). [Paper]
CVPR 2025 Workshop on Responsible Generative AI (ReGenAI '25)

2024 "Why Am I Still Seeing This: Measuring the Effectiveness of Ad Controls and Explanations in AI-Mediated Ad Targeting Systems"
Jane Castleman and Aleksandra Korolova, 2024.
7th AAAI Conference on AI, Ethics, and Society (AIES '24). [Paper]
Fourth European Workshop on Algorithmic Fairness (EWAF '25).

Workshop papers

2025 "Missing the Mark: Rethinking Math Benchmarks for LLMs using IRT"
Jane Castleman*, Nimra Nadeem*, Tanvi Namjoshi*, and Lydia T. Liu, 2025.
AAAI 2025 Workshop on AI for Education (iRAISE). **(Spotlight)** [Paper] **denotes equal contribution.*

Invited Talks

2026 Auditing Proprietary Platforms Through the Private Aggregation of Individual Reports
Data-Driven Social Science Connect: Interdisciplinary Exchange

2026 Third-Party Evaluations: Methods and Implications
CITP State AI Policy Forum

2025 Rethinking Benchmarks: Designing Valid & Reliable Evaluations for Generative AI
Data-Driven Social Science Connect: Interdisciplinary Exchange

2024 Ad Controls & Explanations with AI-mediated Ad Targeting
CITP Conference Tech Policy: The Next 10 Years

Grants, honors & awards

2026 Siebel Scholar (One of 5 Princeton computer science graduate students selected for an award of $35,000)
Graduate Fellowship in Social Data Science (One of 9 Princeton graduate students selected)
2025 Princeton Data-Driven Social Science Grant ($9,400)
Heidelberg Laureate Forum Young Researcher (One of 100 computer science students chosen globally)
2024 Outstanding Student Teaching Award, Princeton CS Dept. ($1,000)
McIntosh Senior Thesis Funding Award, Princeton School of Engineering ($5,000)
2023 All-Academic Team (National Intercollegiate Rugby Association)

Research & Professional Experience

2024-2026 **Center for Information Technology Policy**
Research Assistant, Advisor: Aleksandra Korolova

Initiating and leading multiple concurrent projects to publication, studying user privacy on generative AI platforms, limitations of model alignment, and evaluations of bias in generative AI.

2023 **Ida B. Wells Just Data Lab**
Research Associate, Advisor: Ruha Benjamin

Led analysis of employment barriers for justice-impacted individuals across 41 U.S. occupations; quantified monetary losses ($25k–$250k) due to occupational licensing restrictions.

Teaching

2026 **COS351: Technology Policy & Law**
Teaching Assistant

Organizing discussion sections and advising student semester-long research projects relating to AI governance, auditing, data privacy, and cybersecurity.

2025 **COS350: Ethics of Computing**
Head Teaching Assistant

Leading group of 4 TAs to plan precepts, help design assignments, & grade assignments. Hold weekly office hours and/or precepts explaining fundamental course concepts, such as fairness in ML, bias in AI, and differential privacy. Creating new teaching materials explaining these concepts to students.